

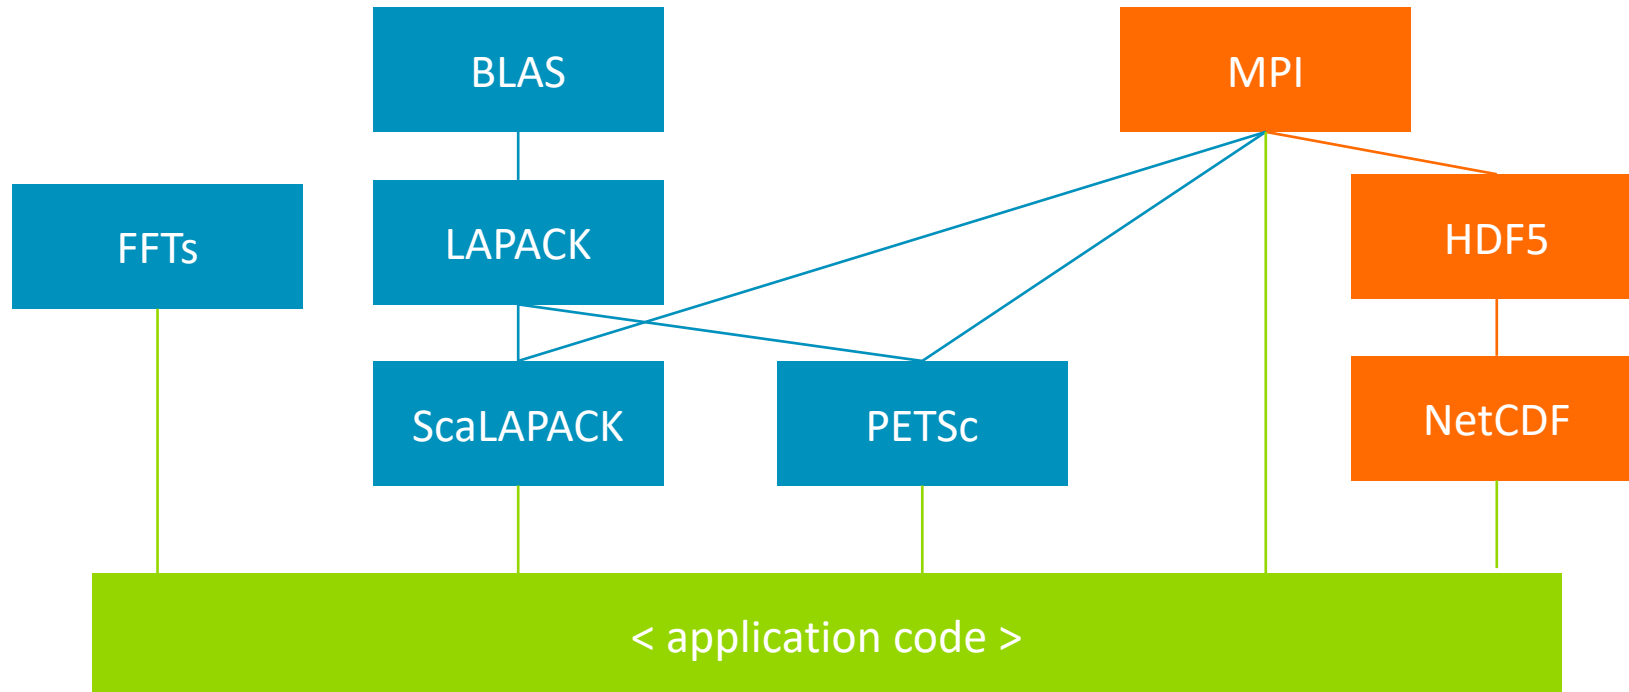


arm

A brief introduction to open source maths libraries & Arm Performance Libraries

Maths libraries on Arm for HPC

- Numerical libraries are the bedrock of most scientific codes solved on HPC systems
 - Re-inventing the wheel is normally a bad idea
- There are hierarchies of tools that are built on each other
 - Today many of these work out-of-the-box
 - Those that need specific Arm optimizations are already starting to get these added
 - There are a mix of open source and proprietary libraries



Arm Performance Libraries

Optimized BLAS, LAPACK and FFT



Commercially supported
by Arm



Best in class performance



Validated with
NAG test suite

Commercial 64-bit Armv8-A math libraries

- Commonly used low-level math routines - BLAS, LAPACK and FFT
- Provides FFTW compatible interface for FFT routines
- Sparse linear algebra and batched BLAS support
- libamath gives high-performing implementations of math.h function

Best-in-class serial and parallel performance

- Algorithmic, Armv8-A and SVE optimizations
- Tuning for specific platforms like Neoverse N1 in AWS

Validated and supported by Arm

- Validated with NAG's test suite, a de-facto standard
- Available in both commercial and free forms

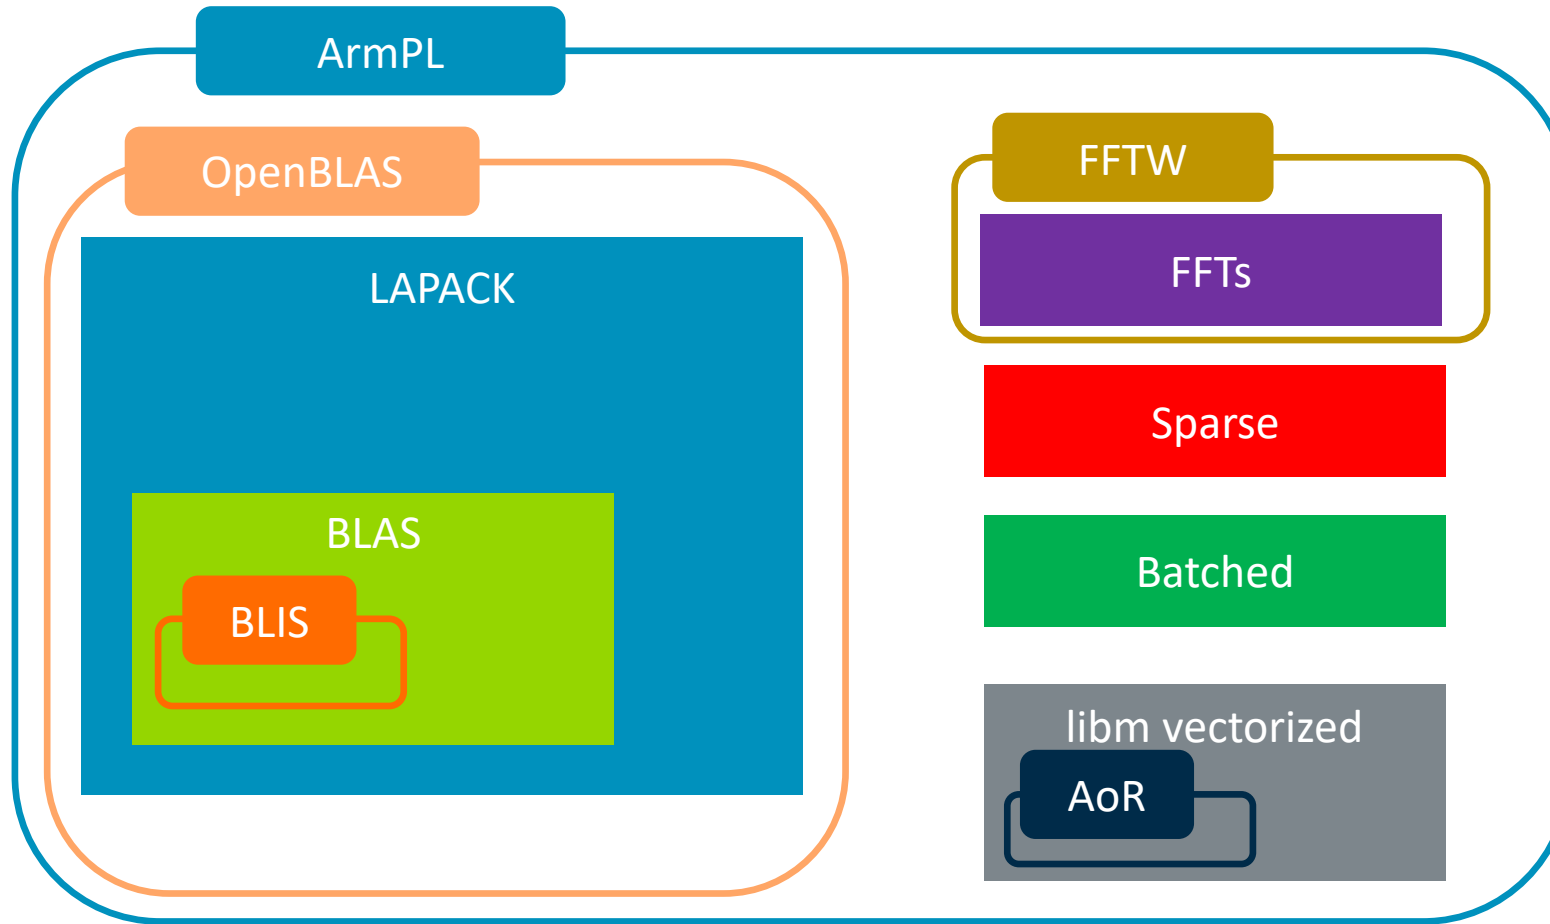
What does all this mean in practice?

Commercial 64-bit Armv8 math libraries

- These libraries are provided both as part of the paid-for *Arm Compiler for Linux* product and as a free version
 - ACfL
 - Versions compatible with both GCC and Arm Compiler for Linux
 - Microarchitectural optimizations included for Arm partner cores as well as Cortex/Neoverse ones
 - Free version
 - Only compatible with GCC
 - Microarchitectural optimizations included for only Arm-designed cores, such as Neoverse N1, as found in the AWS Graviton2 systems
- We are what is known as “the vendor maths library”
 - This means we should generally be the end user’s best option for the highest performing implementations of the functionality of the architecture
 - Open source alternatives may beat performance in some cases on some platforms
 - Other examples are Intel MKL (on x86) and IBM’s ESSL (on POWER)
 - Some systems integrators will also provide their own, e.g. Cray’s libsci

What does Arm Performance Libraries provide?

...and how is this different to open source libraries



Key: Blocks are technologies
 Curved shapes are libs

Notes:

- **OpenBLAS** mainly only optimizes BLAS
- **BLIS** provides excellent optimizations for a subset of BLAS functions
- **ArmPL** is also a delivery vehicle for *Arm Optimized Routines*, but also develop implementations to be open-sourced

Choices in which library version users want

- Unfortunately, there are a few fundamental choices that users can make which mean that we cannot just ship a single library
 - **OpenMP**
 - Users may want the maths library they are using to take advantage of all cores available, or they may not want any OpenMP done by our library.
 - Nested parallelism may be desired by the user
 - **Having 32-bit or 64-bit integers (lp64 or ilp64)**
 - This issue confuses everyone at some stage. If you need to run a code with very large counts of *something* then it is necessary to use 64-bit integers
 - **Compiler choice**
 - If using the full ACfL implementation, users must also choose to have the library match a GCC or an Arm Compiler build
 - Arm Compiler has added integration options that rather than having to explicitly specify the include directory and then link the library the “-armpl” option can be used in compile and link stages
- Open source libraries typically offer these same choices
 - They will be build-time options so you need to compile appropriately
- For **FFTW** the choice of single or double precision is also a buildtime choice

Documentation

- We have full documentation online at
 - <https://developer.arm.com/documentation/101004/latest>
- This includes full descriptions of the APIs we support
 - This includes FFTW calls!
 - Google for, e.g., “Arm DGEMM” normally gets our page as the first hit
- We also ship examples that give working cases illustrating how to call, compile and link
 - Both C and Fortran examples provided
- There is also a Getting Started guide online
 - <https://developer.arm.com/tools-and-software/server-and-hpc/downloads/arm-performance-libraries/get-started-with-armpl-free-version>

The screenshot shows the 'Arm Developer' website with the 'Arm Performance Libraries Reference Guide' for 'fftw_plan_dft'. The page layout includes a top navigation bar, a search bar, and a table of contents on the left. The main content area provides a detailed description of the 'fftw_plan_dft' function, including its purpose, a note about 'in' and 'out' pointers, and syntax for both C and Fortran. The C code example shows the inclusion of 'fftw3.h' and the use of 'fftw_plan_dft' to create a plan. The Fortran code example shows the use of 'fftw_plan_dft' to create a plan. The 'Returns' section states that an 'fftw_plan' object is returned. The 'Parameters' section lists 'rank' as an input parameter, which is the number of dimensions of the FFT problem to be solved, with a constraint that 'rank' must be greater than or equal to 1.



arm

BLAS and LAPACK

Commonly used low-level math routines

- The standard linear algebra libraries used in HPC are **BLAS** and **LAPACK**
- Most routines come in a four varieties (where appropriate)
 - Single precision real : Routines prefixed by 'S'
 - Double precision real : Routines prefixed by 'D'
 - Single precision complex : Routines prefixed by 'C'
 - Double precision complex : Routines prefixed by 'Z'
- The rest of the name (normally) describes something about what the routine does
 - E.g. the matrix-matrix multiplication routine **DGEMM** is a
 - **D** – Double precision
 - **GE** – Matrices given in General format
 - **MM** – Matrix-Matrix multiplication is performed

Linear algebra

BLAS

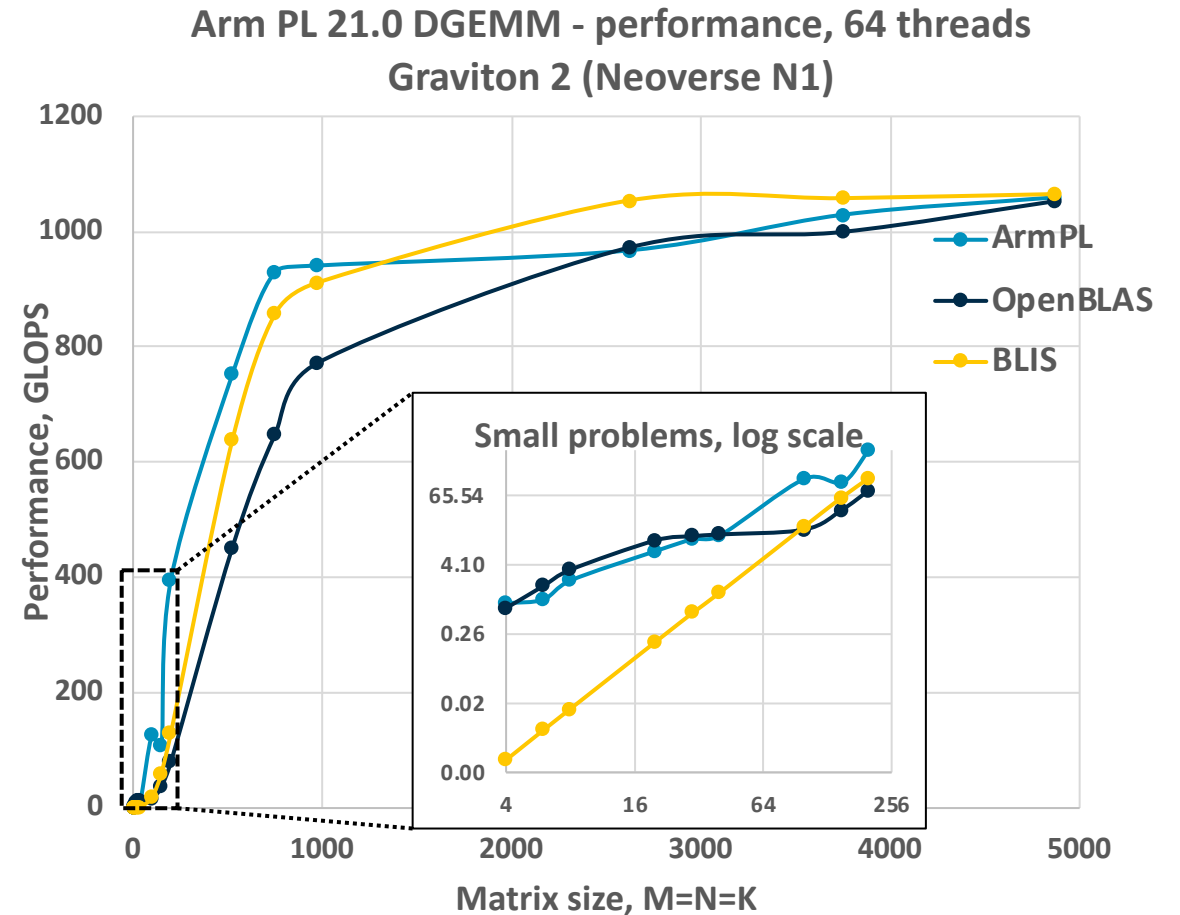
- Basic Linear Algebra Subroutines, is a standard API
 - It is provided on all systems, used by scientific codes for vector and matrix maths
 - It was designed for Fortran, but is callable from all languages
- These routines are come in three levels
 - BLAS level 1 – vector-vector operations,
 - e.g. DCOPY, DAXPY, DDOT
 - BLAS level 2 – matrix-vector operations,
 - e.g. DGEMV, DTRMV, DGER
 - BLAS level 3 – matrix-matrix operations,
 - e.g. DGEMM, DTRMM, DTRSM
- 156 BLAS routines in total

LAPACK

- [LAPACK](#), the Linear Algebra Package, is another standard API
 - It is provided on all systems, used by a wealth of scientific codes for solving equation systems
 - It was designed for Fortran, but is callable from all languages
 - We currently support LAPACK 3.9.0
- LAPACK is built on BLAS routines
- There are now around 1700 LAPACK routines


Micro-architectural tuning

- For the best performance possible micro-architectural tuning is needed
- All BLAS kernels are handwritten in *assembly code* in order to maximise overall performance
- Different micro-architectures may need differences in the instruction ordering – or even the instructions used
- In **ArmPL** the choice of implementation to be used is done automatically at runtime based on the hardware used
 - Open source libraries do not do this



Arm Performance Libraries: OpenMP Scaling on N1

Run on AWS Graviton2



- Shown is DGEMM on square matrices using 64 threads on an AWS Graviton2
- Shown for matrix sizes of 100, 1,000 and 10,000
- Shows up to over 90% efficiency for large matrices at 64 cores



arm

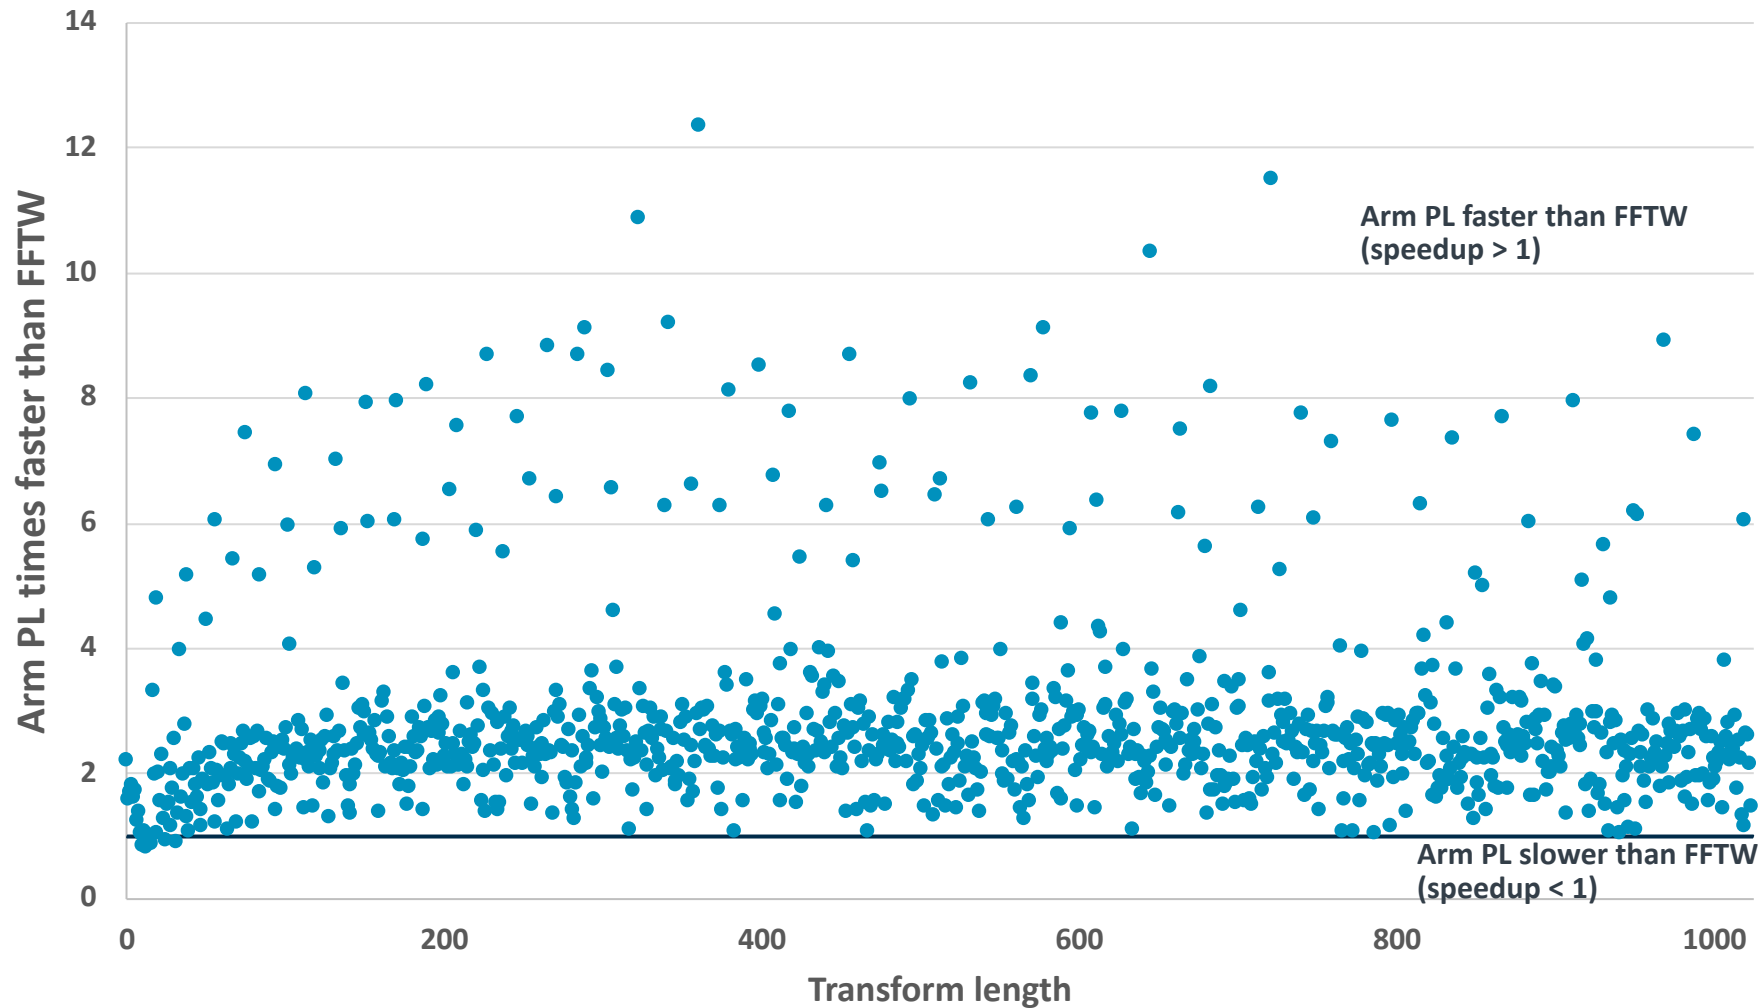
FFTs

Fast Fourier Transforms

- FFTs are very commonly used in a wide variety of applications. They allow some hard problems to be transformed into a way that can be solved much more easily.
- FFTs are solved by taking time to do a “planning” stage first, before doing many fast “execute” calls of that plan
- The open-source **FFTW3** interface has become the de facto standard in scientific applications
- **Arm Performance Libraries** includes the full set of all the Discrete Fourier Transform (DFT) functions using the **FFTW3** interface
 - All functions available in:
 - 1-d, 2-d, 3-d and n-d
 - Complex-to-complex, real-to-complex and complex-to-real
 - All Discrete Cosine Transform (real-to-real) functions also supported
 - Single, double **and half precision**
 - FFTW’s MPI interface is also supported
 - All FFTs transparently generated using a JIT compiler in the library

ArmPL 21.0 FFT vs FFTW 3.3.9


Complex-to-complex single precision 1-d transforms
Graviton 2 (Neoverse N1)



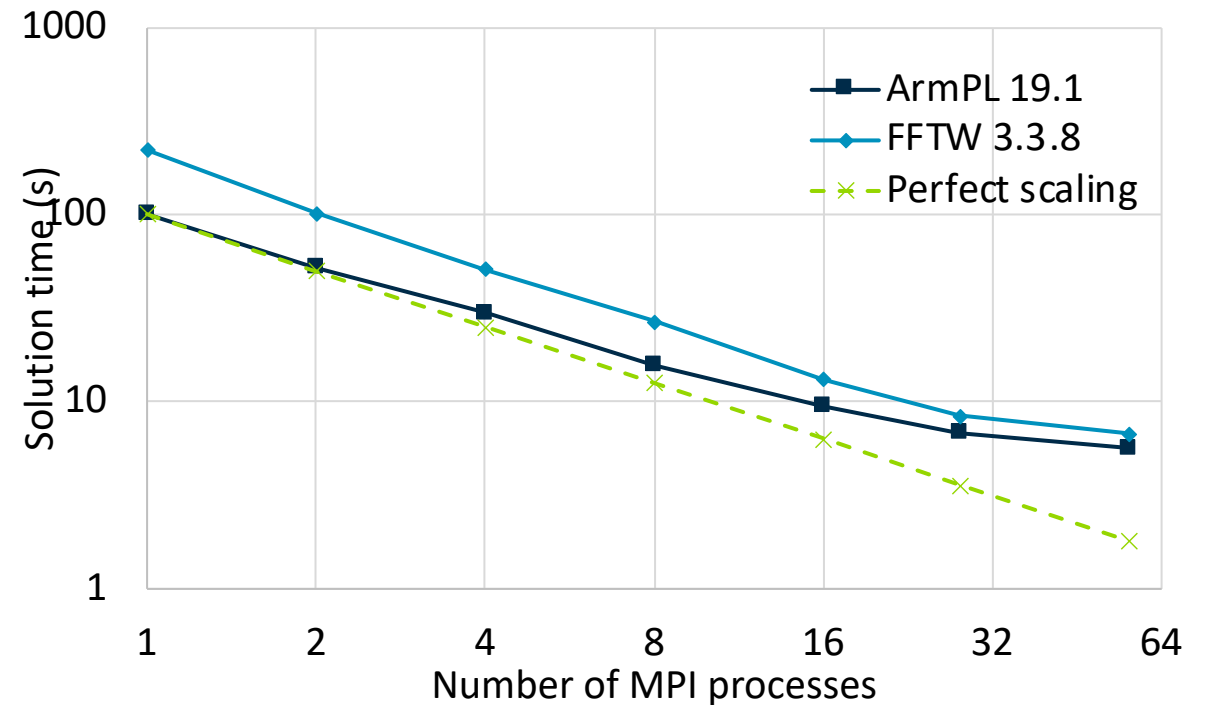
FFT MPI performance – 19.1

Scaling using FFTW MPI interface improved; now similar scaling to FFTW

FFT MPI performance on ThunderX2
3-d case: 512x512x512



FFT MPI performance on ThunderX2
3-d case: 1024x1024x1024






arm

Sparse matrix
support

Sparse matrix support

- Vendor maths libraries have not standard (like LAPACK) to use for high performing implementations of sparse functions
- **Arm Performance Libraries** has added high performing implementations of SpMV and SpMM
 - These rely on the code writer knowing about the possibility of reuse of matrices which is where the optimization potential (over CSR) comes from
- Operation schematic:
 - Create and fill sparse matrix object
 - Optimize for operation
 - Call <operation> many times

Speed-up for SpMV compared to CSR on a selection of sparse matrices





arm

Batched BLAS/LAPACK interfaces

Batched BLAS


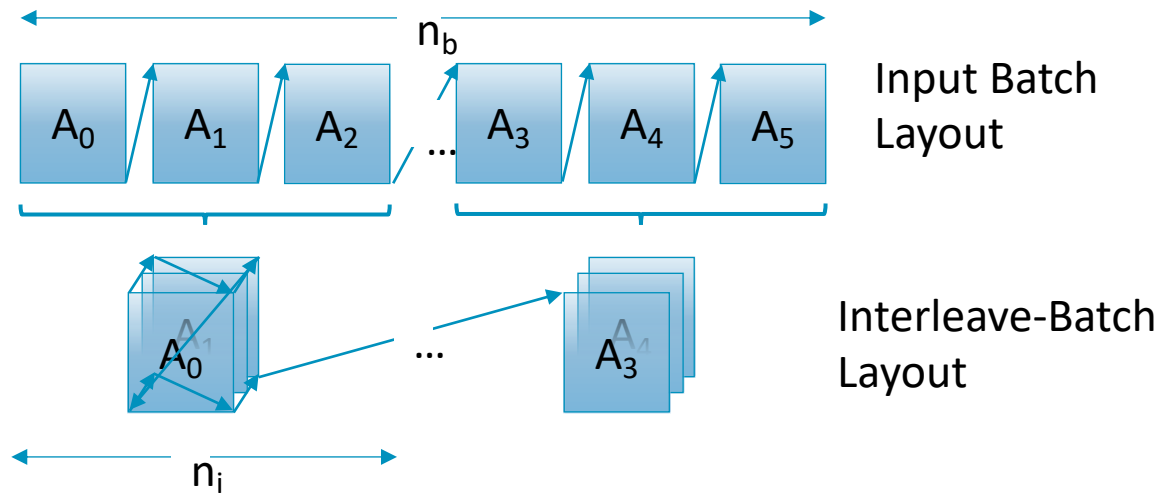
- Batched linear algebra is a growing area of importance for many real scientific codes
- This enables users to make a single library call and ask for the same operation, e.g. DGEMM, to be done on an array of input matrices
- This has the advantage of missing out the overhead of input parameter checking (which can be substantial on tiny cases) and allowing the library to do optimizations that would not be possible in individual calls
 - Parallelisation over sets of cases is an easy example of how this can be done
- **Arm Performance Libraries** support a batched operations using `*gemm_batched()`
Even higher performance is possible if you know that data will be reused more than once, or you have more control of your input data
 - For these cases ArmPL has developed a set of functions where the data gets pre-interleaved...

New in ArmPL 21.0: Interleave Batch layout

Interleaved batched BLAS/LAPACK functions

Using batched interfaces

- Executing a routine on a batch of n_p matrices
- Split batch of matrices into n_b times length- n_i (sub-)batches



Functions supported

- BLAS: ddot, dger, dgemm, dgemv, dscal, dtrmm, dtrsm, dtrsv
- LAPACK: dgeqrf (QR), dgetrf (LU), dpotrf (Cholesky)
 - plus dormqr, dorgqr for multiplying and generating Q)




arm

libamath

Optimized libm functions

Open Source: <https://github.com/ARM-software/optimized-routines>

Normalised runtime



ArmPL includes libamath and libastring

- Algorithmically better performance than standard library calls
- No loss of accuracy
- Enabled by default with Arm Compiler for Linux
- Single and double precision implementations of:
`erf()`, `erfc()`
- single and double precision implementations of:
`exp()`, `pow()`, `log()`, `log10()`
- single precision implementations of:
`sin()`, `cos()`, `sincos()`
- Efficient **memory/string** functions from `string.h`
- Enable **autovectorization** of math and string routines in ACfL by adding `-armpl` or `-fsimdmath`

arm

How to use – including
with SPACK

Code changes needed to use Arm Performance Libraries

- If your code is already set up to use another BLAS, LAPACK or FFTW3 implementation then it will just work!
 - All functions use the standard APIs
 - The only exceptions are the sparse and batched interfaces
- If using the FFTW MPI interface you must ensure that the ArmPL version of “fftw3.h” is used rather than the FFTW version
- The C correct header file to include is “armpl.h”
 - although “blas.h” and “lapack.h” will also work
- In Fortran 90 and beyond you may need to use the module “armpl_library”

Arm Performance Libraries – linking

- In order to compile applications using BLAS, LAPACK and FFT routines from the Arm Performance libraries, four options are provided:
 - Serial and OpenMP builds
 - 32-bit and 64-bit integers
- These translate into four static binaries in `/opt/arm/armpl-*/lib/`
 - `libarmpl_lp64.a`
 - `libarmpl_lp64_mp.a`
 - `libarmpl_ilp64.a`
 - `libarmpl_ilp64_mp.a`
- Shared libraries (`libarmpl*.so`) are also provided
- Compile and link using, for example

```
gcc -O3 file.c -fopenmp -c file.o -I${ARMPL_DIR}/include
gcc -O3 file.o -fopenmp -o file -L${ARMPL_DIR}/lib -larmpl_mp
```

- `libamath` will be automatically linked in
 - Vector calls to maths functions only from Arm Compiler for Linux
 - GCC cannot automatically generate them at this time.

Spack virtual packages

- Spack provides the concept of a virtual package
 - A package depends upon an implementation, but doesn't necessarily care whose
 - Used for MPI, BLAS, LAPACK, FFTW-API
- Applications depend on the virtual package
- Different packages can provide those implementations
 - Choice is made at installation time – user can override

App Using Virtual Package (SW4LITE)	Library Providing Implementation (ArmPL)	Library Providing Implementation (Open MPI)
depends_on('blas')	provides('blas')	provides('mpi')
depends_on('lapack')	provides('lapack')	provides('mpi@:3.1', when='@2.0.0:')
depends_on('mpi')	provides('fftw-api@3')	

Spack specs for virtual packages

Ask Spack to fulfil with defaults: Results in ArmPL and OpenMPI

```
$ spack spec sw4lite%gcc
...
sw4lite@1.1%gcc@10.3.0~ckernel+openmp precision=double arch=linux-amzn2-graviton2
  ^armpl@21.0.0%gcc@10.3.0~ilp64~shared threads=none arch=linux-amzn2-graviton2
  ^openmpi@4.1.0%gcc@10.3.0~atomics~cuda~cxx~cxx_exceptions .....
```

Ask Spack to fulfil with 'openblas': Results in OpenBLAS and OpenMPI

```
$ spack spec sw4lite%gcc ^openblas
...
sw4lite@1.1%gcc@10.3.0~ckernel+openmp precision=double arch=linux-amzn2-graviton2
  ^openblas@0.3.15%gcc@10.3.0~bignuma~consistent_fpcsr~ilp64+locking+pic+shared
    threads=none arch=linux-amzn2-graviton2
  ^openmpi@4.1.0%gcc@10.3.0~atomics~cuda~cxx~cxx_exceptions .....
```

Who provides the implementations?

BLAS	LAPACK	FFTW-API
armpl	armpl	amdfftw
atlas	atlas	armpl
blis	cray-libsci	cray-fftw
cray-libsci	flexiblas	fftw
essl	fujitsu-ssl2	fujitsu-fftw
fujitsu-ssl2	intel-mkl	intel-mkl
intel-mkl	intel-oneapi-mkl	intel-oneapi-mkl
intel-oneapi-mkl	intel-parallel-studio	intel-parallel-studio
intel-parallel-studio	libflame	
netlib-lapack	netlib-lapack	
netlib-xblas	nvhpc	
nvhpc	openblas	
openblas	veclibfort	
veclibfort		

How to use the virtual packages?

- Each implementation provides a 'libs'
 - Listing which libraries to use, and their paths
 - Wrapped up as 'blas_libs', 'lapack_libs' and 'fftw_libs'
- Using package can then just reference the virtual package:

```
lapack_blas = spec['lapack'].libs + spec['blas'].libs
targets.append('EXTRA_LINK_FLAGS={0}'.format(lapack_blas.ld_flags))
```

- Results in: (for sw4lite%gcc ^armpl)

```
EXTRA_LINK_FLAGS= \
-L/software/ACFL/21.0/armpl-21.0.0_AArch64_RHEL-7_gcc_aarch64-linux/lib \
-L/lib64 \
-L/software/gcc/10.3.0/lib64 \
-larmpl -lamath -lastring -lm -lgfortran -lgfortran
```

How are variants handled? ... It's complicated

- Variants
 - OpenMP support ('threads=openmp'), 64-bit int support ('+ilp64')
 - Normally get passed through as a dependency
- Different packages handle variants differently
 - So, we need to specify them explicitly

```
depends_on('blas')  
depends_on('lapack')  
depends_on('armpl+ilp64 threads=openmp', when='^armpl')  
depends_on('openblas+ilp64 threads=openmp', when='^openblas')
```


A man with glasses and a grey t-shirt is sitting on a wooden coffee table, working on a laptop. He is looking at the screen and typing. In the background, two children are playing on the floor with toys. The room is a living room with a red armchair and a white fireplace. The word 'arm' is overlaid in white text on the left side of the image.

arm


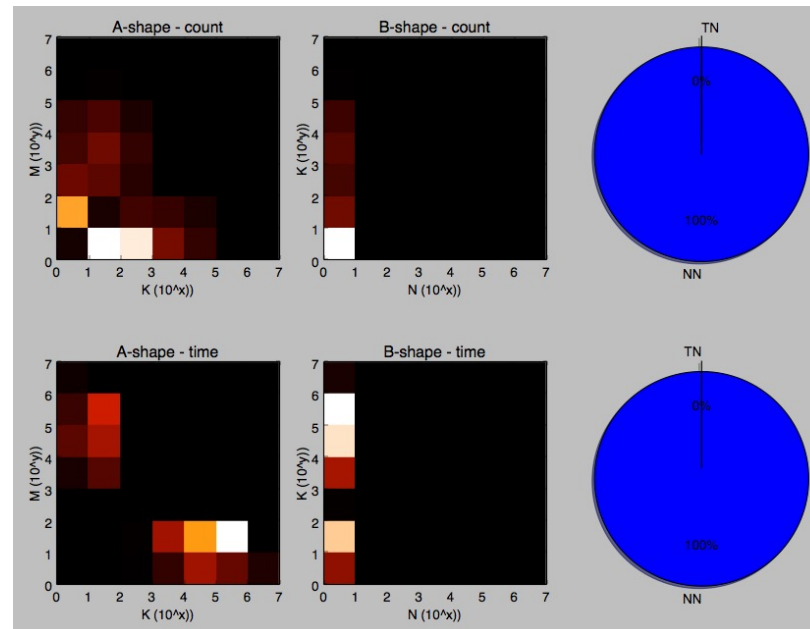
Logging library calls: perf-libs-tools

Logging Library : perf-lib-tools

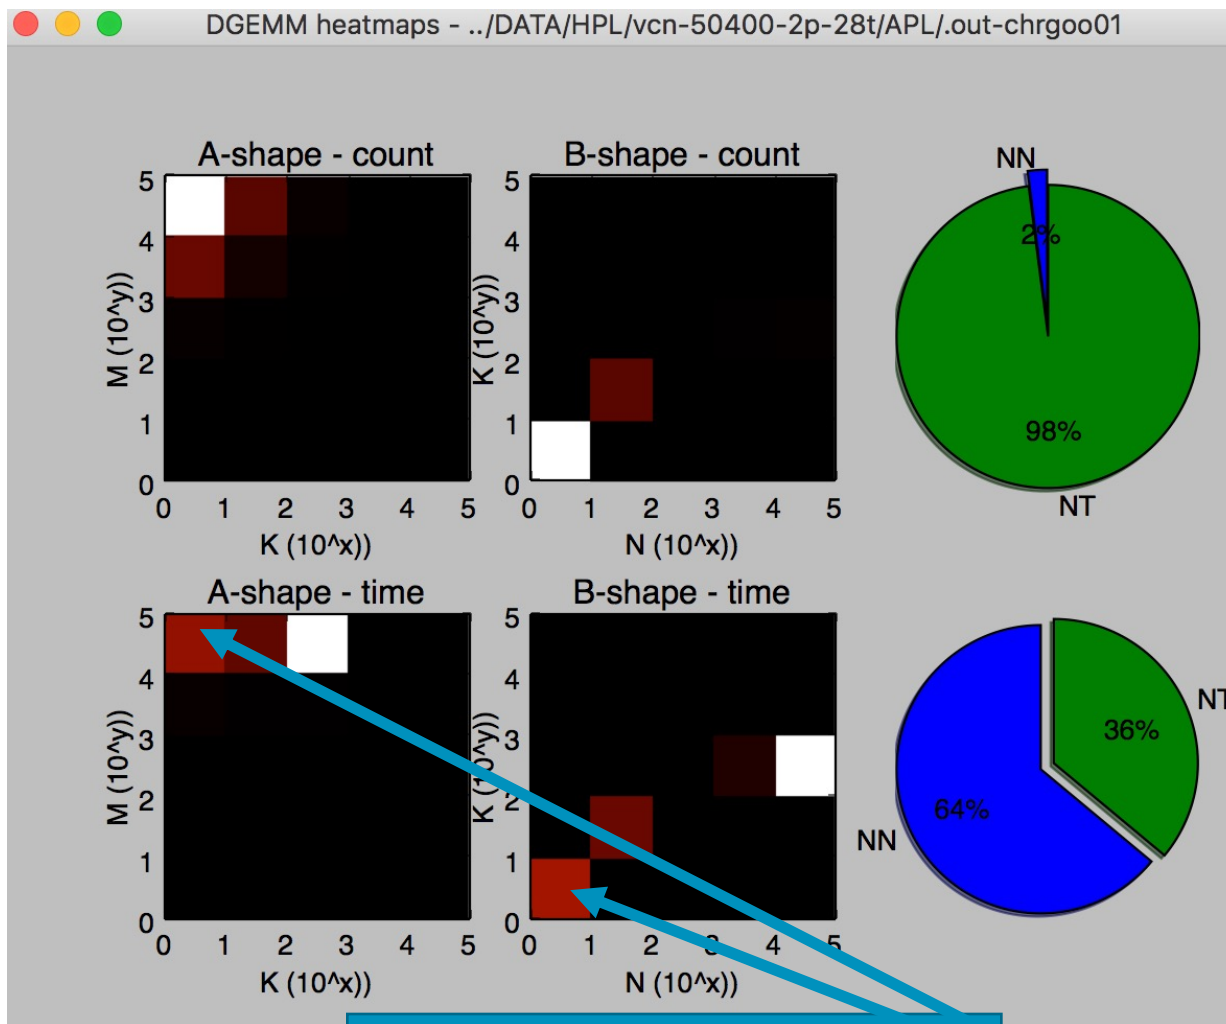
<https://github.com/ARM-software/perf-lib-tools>

Understanding an application's needs for BLAS, LAPACK and FFT calls

- Used in conjunction with **Arm Performance Libraries** can generate logging info to help profile applications for specific case breakdowns
- Allows us to identify where time is spent in real applications.

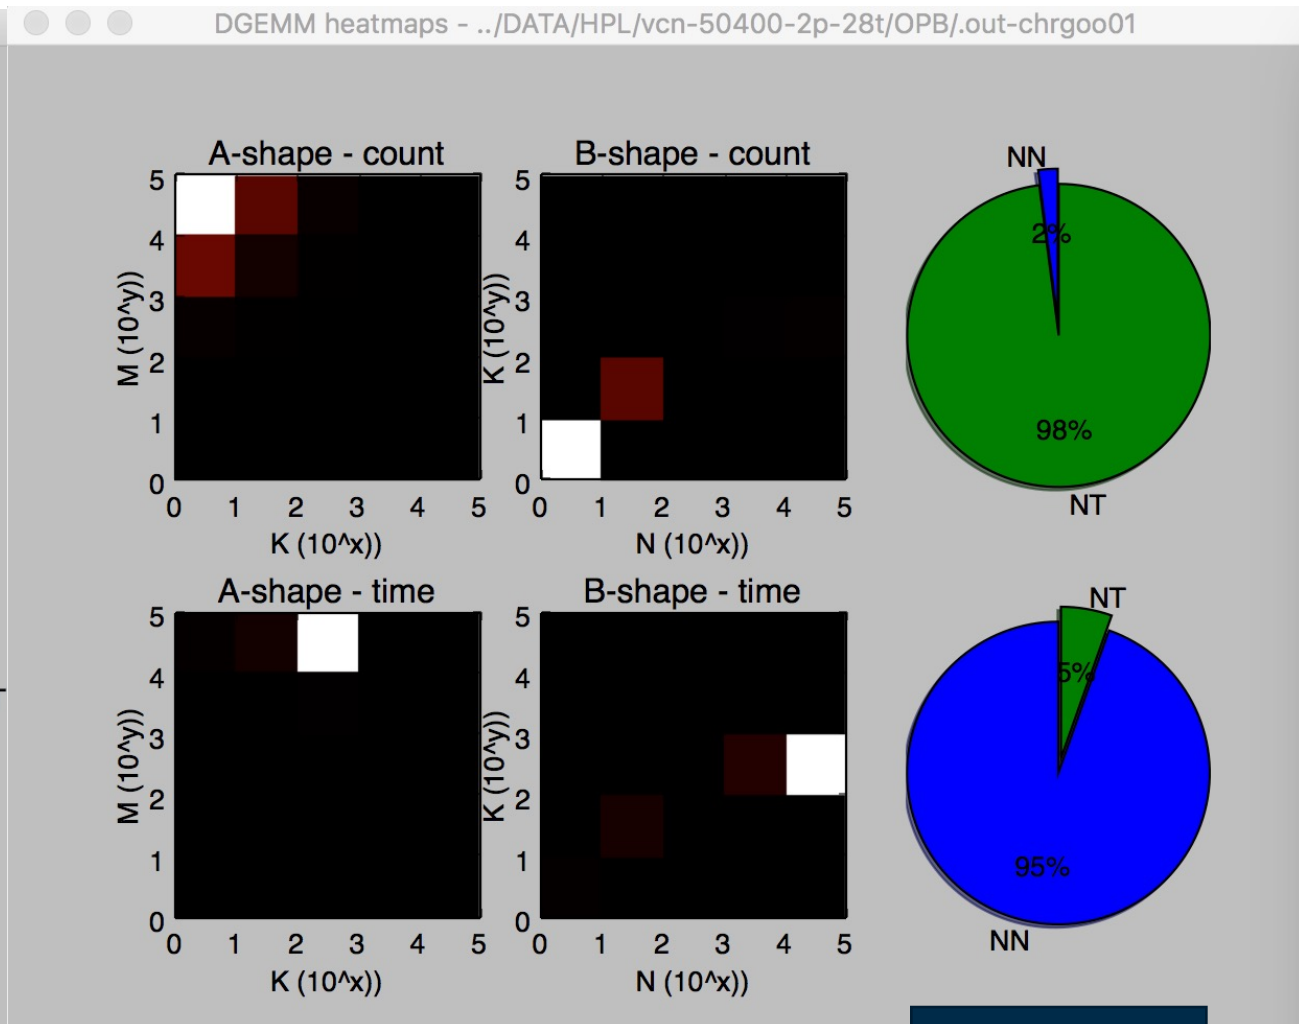


Ages ago OpenBLAS did HPL better. Why?



Arm Performance Libraries

Huge time overhead for tiny cases. Now fixed!



OpenBLAS

arm

perf-libs-tools – Usage guide

Building and collecting data is very straightforward:

```
# git clone https://github.com/ARM-software/perf-libs-tools
# cd perf-libs-tools
# make
# export LD_PRELOAD=$PWD/lib/libarmpl_summarylog.so
# ./application
<output>
Arm Performance Libraries output summary stored in
/tmp/armplsummary_16776.apl
```

Using perf-libs-tools then running the “process_summary.py” script gives lots of useful information like:

```
# ~/perf-libs-tools/tools/process_summary.py /tmp/armplsummary_*
Process full dataset for BLAS, LAPACK and FFT function usage.
...
BLAS level 1      : count    88218776    total time    48.6955
BLAS level 2      : count   334145452    total time    1672.3763
BLAS level 3      : count    7473286    total time    151.5514
LAPACK            : count    325027     total time     5.2572
FFT               : count   60446086    total time    481.7579
```

perf-lib-tools: Summary output per function

BLAS cases:

BLAS level 1:

dcopy_	cnt=	17440189	totTime=	24.0269
dscal_	cnt=	67608016	totTime=	20.9392
zcopy_	cnt=	2232038	totTime=	3.0867
zdotc_	cnt=	175374	totTime=	0.2762
zdscal_	cnt=	322549	totTime=	0.2116
zscal_	cnt=	230973	totTime=	0.0696
zaxpy_	cnt=	50006	totTime=	0.0201
idamax_	cnt=	48546	totTime=	0.0188
zswap_	cnt=	39103	totTime=	0.0116
dznrm2_	cnt=	15485	totTime=	0.0097
ddot_	cnt=	14859	totTime=	0.0064
dzasum_	cnt=	15924	totTime=	0.0056
daxpy_	cnt=	14859	totTime=	0.0055
dasum_	cnt=	6216	totTime=	0.0031
dnrm2_	cnt=	2072	totTime=	0.0021
izamax_	cnt=	1067	totTime=	0.0017
dswap_	cnt=	1500	totTime=	0.0006

BLAS level 2:

dgemv_	cnt=	333911120	totTime=	1671.7741
zgemv_	cnt=	137302	totTime=	0.4562
ztrmv_	cnt=	27594	totTime=	0.0466
ztrsv_	cnt=	23587	totTime=	0.0315
zhemv_	cnt=	13104	totTime=	0.0292
zher2_	cnt=	15081	totTime=	0.0279
dtrsv_	cnt=	11840	totTime=	0.0072
zgerc_	cnt=	5824	totTime=	0.0036

BLAS level 3:

dgemm_	cnt=	7150124	totTime=	90.5322
zgemm_	cnt=	201115	totTime=	59.9991
ztrmm_	cnt=	3562	totTime=	0.7432
ztrsm_	cnt=	42949	totTime=	0.1444
zherk_	cnt=	43304	totTime=	0.0813
dtrsm_	cnt=	31008	totTime=	0.0358
zher2k_	cnt=	928	totTime=	0.0103
zhemm_	cnt=	296	totTime=	0.0050

LAPACK cases:

zhegv_	cnt=	843	totTime=	1.2420
zheev_	cnt=	843	totTime=	0.7431
zpotrf_	cnt=	1750	totTime=	0.4506
dgetrf_	cnt=	5920	totTime=	0.4397
zhegst_	cnt=	843	totTime=	0.3478
zhegs2_	cnt=	991	totTime=	0.3318
zsteqr_	cnt=	843	totTime=	0.2665
zhetr_	cnt=	843	totTime=	0.2627
zungtr_	cnt=	843	totTime=	0.1859
ztrtri_	cnt=	64	totTime=	0.1818
zpocon_	cnt=	843	totTime=	0.1629

Case Study: (1) Report from a user

A user reported that the 18.4 libraries release was slower than FFTW for a key application

Compiler	BLAS	FFT	Runtime (s)
GCC 7.2	ArmPL 18.4	Cray FFTW	195
GCC 7.2	OpenBLAS 0.2	Cray FFTW	210
GCC 7.2	ArmPL 18.4	ArmPL 18.4	237
GCC 7.2	OpenBLAS 0.2	ArmPL 18.4	252

We were confused since we were confident that we could solve the sized FFT cases encountered quicker...

Case Study: (2) Using perf-libcs-tools

The user was then able to provide full perf-libcs-tools for the whole solution run.

The FFT summaries looked like:

FFT cases:

FFTW calls:

fftw_plan_many_dft	len= [54]	plan-cnt=	20092	plan-Time=	5.735552	exec-cnt=	24126	exec-Time=	20.338287
fftw_plan_many_dft	len= [60]	plan-cnt=	20151	plan-Time=	5.380509	exec-cnt=	18043	exec-Time=	15.108546
fftw_plan_many_dft	len= [80]	plan-cnt=	20151	plan-Time=	2.144119	exec-cnt=	18262	exec-Time=	4.805128
fftw_plan_many_dft	len= [120]	plan-cnt=	404	plan-Time=	0.038137	exec-cnt=	444	exec-Time=	0.081618
fftw_plan_many_dft	len= [160]	plan-cnt=	404	plan-Time=	0.058282	exec-cnt=	465	exec-Time=	0.115087
fftw_plan_many_dft_c2r	len= [54]	plan-cnt=	30	plan-Time=	0.001394	exec-cnt=	0	exec-Time=	0.000000
fftw_plan_many_dft_c2r	len= [108]	plan-cnt=	188	plan-Time=	0.042434	exec-cnt=	139	exec-Time=	0.041167
fftw_plan_many_dft_r2c	len= [54]	plan-cnt=	29	plan-Time=	0.001645	exec-cnt=	0	exec-Time=	0.000000
fftw_plan_many_dft_r2c	len= [108]	plan-cnt=	216	plan-Time=	0.043146	exec-cnt=	121	exec-Time=	0.027025

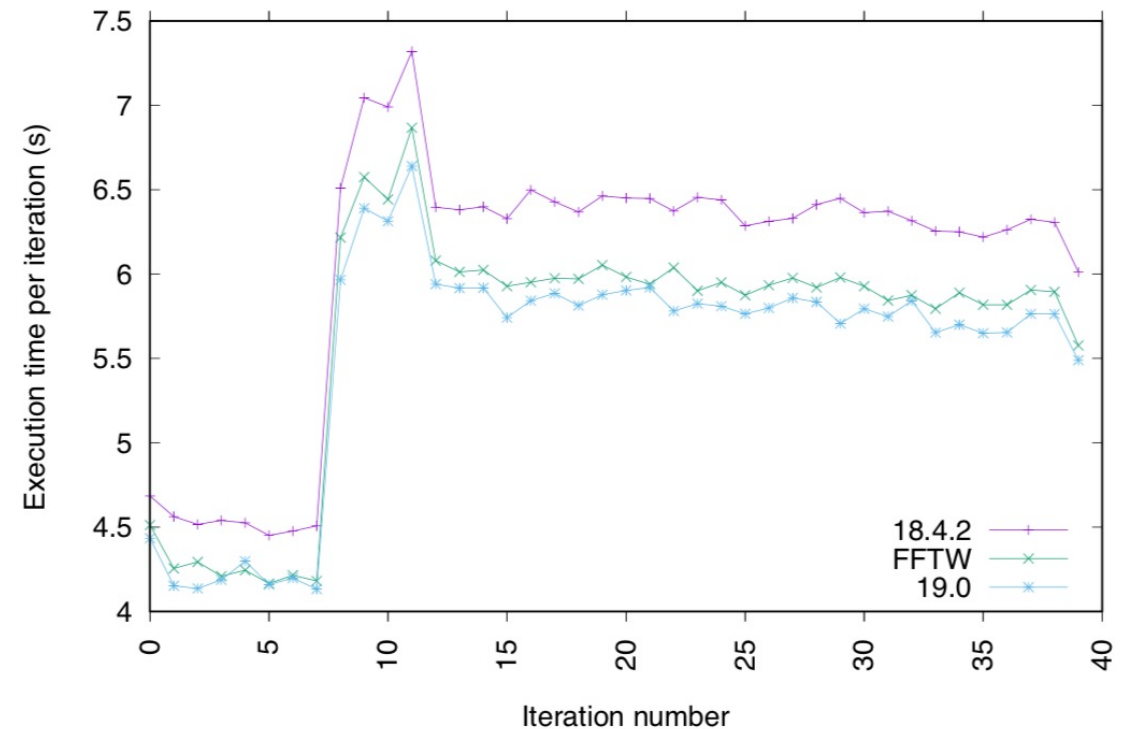
Library	Planning time	Execution time
ArmPL 18.4.0	13.66s	41.24s
FFTW	4.82s	49.72s

Case Study: (3) Fixing the issue

Back to the innards of the libraries

- It was clear that the code was creating one FFT plan per execution – a use-case we had never expected
- When we plan our FFTs various data-structures are set up.
- This includes allocating certain working arrays, even if we don't need to do auditioning as the case has been put into 'wisdom' previously
- It turned out we were unnecessarily zeroing an array which meant that if planning wasn't a rare event, there was an overhead
- Fixing this was very quick to implement and in the released 19.0 a week later

Final performance





Getting more information

We want you to be as successful as possible!

- If you have a question about anything to do with linking or running the various options do ask on the '*#help-maths-libraries*' channel on the A-HUG Cloud Hackathon Slack
 - We've got lots of experience of these packages. Don't suffer in silence!
- ArmPL:
 - Documentation:
 - <https://developer.arm.com/documentation/101004/latest>
 - Getting Started guide online
 - <https://developer.arm.com/tools-and-software/server-and-hpc/downloads/arm-performance-libraries/get-started-with-armpl-free-version>
- Perf-libs-tools:
 - <https://github.com/ARM-software/perf-libs-tools>
- OpenBLAS:
 - <https://github.com/xianyi/OpenBLAS>
- BLIS:
 - <https://github.com/flame/blis>
- FFTW
 - <http://www.fftw.org/>

Please do share your perf-libs-tools outputs with us. This will help prioritization of future work.



The Arm trademarks featured in this presentation are registered trademarks or trademarks of Arm Limited (or its subsidiaries) in the US and/or elsewhere. All rights reserved. All other marks featured may be trademarks of their respective owners.

www.arm.com/company/policies/trademarks